

HW6_SVM_Shuting_Li


Shuting

3/09/2022

9.3

(a)

```
x1 = c(3, 2, 4, 1, 2, 4, 4)
x2 = c(4, 2, 4, 4, 1, 3, 1)
colors = rep(c("red", "blue"),c(4,3))
plot(x1, x2, col = colors)
```




Answer:

(b)

Sketch the optimal separating hyperplane, and provide the equation for this hyperplane (of the form (9.1)).

```
p1 <- c(2, (1+2)/2)
p2 <- c(4, (3+4)/2)
a <- ((3+4)/2 - (1+2)/2)/(4-2)
b <- (1+2)/2 - 2*a
plot(x1, x2, col = colors)
abline(b, a)
```



Answer:

(c)

Describe the classification rule for the maximal margin classifier. It should be something along the lines of “Classify to Red if $0 + 1X_1 + 2X_2 > 0$, and classify to Blue otherwise.” Provide the values for 0 , 1 , and 2 .


```
beta0 <- 0.5
beta1 <- -1
beta2 <- 1
```

Answer:

(d)

On your sketch, indicate the margin for the maximal margin hyperplane.

```
plot(x1, x2, col = colors)
abline(b, a)
abline(-1, 1, lty = 2)
abline(0, 1, lty = 2)
```




Answer:

(e)

Indicate the support vectors for the maximal margin classifier.

```
plot(x1, x2, col = colors)
abline(b, a)
arrows(2, 1, 2, 1.5)
arrows(2, 2, 2, 1.5)
arrows(4, 4, 4, 3.5)
arrows(4, 3, 4, 3.5)
```



Answer:


(f) Argue that a slight movement of the seventh observation would not affect the maximal margin hyperplane.

Answer: The seventh observation lies outside of the margin for the maximal margin hyperplane.

(g)

Sketch a hyperplane that is not the optimal separating hyperplane, and provide the equation for this hyperplane.

```
plot(x1, x2, col = colors)
abline(b, 0.8)
```




Answer:

(h)

Draw an additional observation on the plot so that the two classes are no longer separable by a hyperplane.

```
x1 <- append(x1,3.5)
x2 <- append(x2,4)
colors <- append(colors,"blue")
plot(x1,x2,col=colors)
```



Answer:

9.5

We have seen that we can fit an SVM with a non-linear kernel in order to perform classification using a non-linear decision boundary. We will now see that we can also obtain a non-linear decision boundary by performing logistic regression using non-linear transformations of the features.

(a)


```
set.seed(1)
x1 = runif(500) - 0.5
x2 = runif(500) - 0.5
y = 1 * (x1^2 - x2^2 > 0)
```

Answer:

(b)

Plot the observations, colored according to their class labels. Your plot should display X1 on the x-axis, and X2 on the y-axis.

```
color <- rep(NA, length(y))
color[y==0] <- "blue"
color[y==1] <- "red"
plot(x1, x2, col=color)
```



Answer:

(c) Fit a logistic regression model to the data, using X1 and X2 as predictors.

```
fit.logis <- glm(y~x1+x2, data=data.frame(x1,x2,y), family = binomial)
summary(fit.logis)
```

Answer:


```
##
## Call:
## glm(formula = y ~ x1 + x2, family = binomial, data = data.frame(x1,
##      x2, y))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.179  -1.139  -1.112   1.206   1.257
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.087260   0.089579  -0.974   0.330
## x1           0.196199   0.316864   0.619   0.536
## x2          -0.002854   0.305712  -0.009   0.993
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 692.18  on 499  degrees of freedom
## Residual deviance: 691.79  on 497  degrees of freedom
## AIC: 697.79
##
## Number of Fisher Scoring iterations: 3
```

The coefficients are all not significant.

(d)

Apply this model to the training data in order to obtain a predicted class label for each training observation. Plot the observations, colored according to the predicted class labels. The decision boundary should be linear.

```
data=data.frame(x1,x2,y)
pred.fit <- predict(fit.logis, data, type="response")
color.pred <- ifelse(pred.fit>0.5, "red", "blue")
plot(x1,x2,col=color.pred)
```



Answer:

x1

The prediction is far from the true data. The decision boundary is linear.

(e)

Now fit a logistic regression model to the data using non-linear functions of X1 and X2 as predictors (e.g. $X1^2$, $X1 \cdot X2$, $\log(X2)$, and so forth).

```
fit.logis2 = glm(y ~ I(x1^2) + I(x2^2), data = data, family = binomial)
```

Answer:

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(fit.logis2)
```

```
##
```

```
## Call:
```

```
## glm(formula = y ~ I(x1^2) + I(x2^2), family = binomial, data = data)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -0.00667  0.00000  0.00000  0.00000  0.00565
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      7.285     173.960   0.042   0.967
## I(x1^2)      93153.509 860924.357   0.108   0.914
## I(x2^2)     -93265.460 863556.384  -0.108   0.914
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
##      Null deviance: 6.9218e+02  on 499  degrees of freedom
```




```
## Residual deviance: 7.9869e-05 on 497 degrees of freedom
## AIC: 6.0001
##
## Number of Fisher Scoring iterations: 25
```

(f)

Apply this model to the training data in order to obtain a predicted class label for each training observation. Plot the observations, colored according to the predicted class labels. The decision boundary should be obviously non-linear. If it is not, then repeat (a)-(e) until you come up with an example in which the predicted class labels are obviously non-linear.

```
pred.fit2 <- predict(fit.logis2, data, type="response")
color.pred2 <- ifelse(pred.fit2>0.5, "red", "blue")
plot(x1,x2,col=color.pred2)
```




Answer:

(g)

Fit a support vector classifier to the data with X1 and X2 as predictors. Obtain a class prediction for each training observation. Plot the observations, colored according to the predicted class labels.

```
fit.svm = svm(as.factor(y) ~ x1 + x2, data, kernel = "linear", cost = 0.1)
pred.svm = predict(fit.svm, data)
color.svm <- ifelse(pred.svm==1, "red", "blue")
plot(x1,x2,col=color.svm)
```




Answer:

(h)

Fit a SVM using a non-linear kernel to the data. Obtain a class prediction for each training observation. Plot the observations, colored according to the predicted class labels.

```
fit.svm2 = svm(as.factor(y) ~ x1 + x2, data, kernel = "polynomial", degree = 2)
pred.svm2 = predict(fit.svm2, data)
color.svm2 <- ifelse(pred.svm2==1, "red", "blue")
plot(x1,x2,col=color.svm2)
```



Answer:

(i)

Comment on your results.

Answer: It is very convenient to use SVM to find non-linear decision boundary without setting much parameters and predictors.

9.7

In this problem, you will use support vector approaches in order to predict whether a given car gets high or low gas mileage based on the Auto data set.

(a)

Create a binary variable that takes on a 1 for cars with gas mileage above the median, and a 0 for cars with gas mileage below the median.

```
new.var = ifelse(Auto$mpg > median(Auto$mpg), 1, 0)
Auto$mpglevel = as.factor(new.var)
```

Answer:

(b)

Fit a support vector classifier to the data with various values of cost, in order to predict whether a car gets high or low gas mileage. Report the cross-validation errors associated with different values of this parameter. Comment on your results. Note you will need to fit the classifier without the gas mileage variable to produce sensible results.

```
set.seed(1)
tune1 = tune(svm, mpglevel ~ ., data = Auto, kernel = "linear", ranges = list(cost = c(0.01, 0.1, 1, 5,
summary(tune1)
```

Answer:

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.01025641
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-02 0.07653846 0.03617137
## 2 1e-01 0.04596154 0.03378238
## 3 1e+00 0.01025641 0.01792836
## 4 5e+00 0.02051282 0.02648194
## 5 1e+01 0.02051282 0.02648194
## 6 1e+02 0.03076923 0.03151981
```

The best cost is 1.

(c)

Now repeat (b), this time using SVMs with radial and polynomial basis kernels, with different values of gamma and degree and cost. Comment on your results.

```
tune2 = tune(svm, mpglevel ~ ., data = Auto, kernel = "polynomial", ranges = list(cost = c(0.1,
1, 5, 10, 100), degree = c(2, 3, 4)))
summary(tune2)
```

Answer:

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
##   100      2
##
## - best performance: 0.3060897
##
## - Detailed performance results:
##   cost degree   error dispersion
## 1    0.1      2 0.6019231 0.06346118
## 2    1.0      2 0.6019231 0.06346118
## 3    5.0      2 0.6019231 0.06346118
## 4   10.0      2 0.5841667 0.07806609
## 5  100.0      2 0.3060897 0.07318010
## 6    0.1      3 0.6019231 0.06346118
## 7    1.0      3 0.6019231 0.06346118
## 8    5.0      3 0.6019231 0.06346118
## 9   10.0      3 0.6019231 0.06346118
## 10 100.0      3 0.3413462 0.14973473
## 11   0.1      4 0.6019231 0.06346118
## 12   1.0      4 0.6019231 0.06346118
## 13   5.0      4 0.6019231 0.06346118
## 14  10.0      4 0.6019231 0.06346118
## 15 100.0      4 0.6019231 0.06346118
```

```
tune3 = tune(svm, mpglevel ~ ., data = Auto, kernel = "radial", ranges = list(cost = c(0.1,
1, 5, 10, 100), gamma = c(0.01, 0.1, 1, 5, 10, 100)))
summary(tune3)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   10 0.01
##
```

```
## - best performance: 0.02044872
##
## - Detailed performance results:
##      cost gamma      error dispersion
## 1    0.1 1e-02 0.08698718 0.04560056
## 2    1.0 1e-02 0.07160256 0.03373099
## 3    5.0 1e-02 0.05121795 0.02967002
## 4   10.0 1e-02 0.02044872 0.01077927
## 5  100.0 1e-02 0.02044872 0.02020886
## 6    0.1 1e-01 0.07673077 0.03419344
## 7    1.0 1e-01 0.05121795 0.03203768
## 8    5.0 1e-01 0.02557692 0.01709522
## 9   10.0 1e-01 0.02301282 0.02244393
## 10 100.0 1e-01 0.03070513 0.03153205
## 11   0.1 1e+00 0.59461538 0.08083319
## 12   1.0 1e+00 0.06141026 0.03026776
## 13   5.0 1e+00 0.06397436 0.02789391
## 14  10.0 1e+00 0.06397436 0.02789391
## 15 100.0 1e+00 0.06397436 0.02789391
## 16   0.1 5e+00 0.59461538 0.08083319
## 17   1.0 5e+00 0.52051282 0.09421163
## 18   5.0 5e+00 0.51538462 0.10051415
## 19  10.0 5e+00 0.51538462 0.10051415
## 20 100.0 5e+00 0.51538462 0.10051415
## 21   0.1 1e+01 0.59461538 0.08083319
## 22   1.0 1e+01 0.55384615 0.09432787
## 23   5.0 1e+01 0.54358974 0.09085645
## 24  10.0 1e+01 0.54358974 0.09085645
## 25 100.0 1e+01 0.54358974 0.09085645
## 26   0.1 1e+02 0.59461538 0.08083319
## 27   1.0 1e+02 0.59461538 0.08083319
## 28   5.0 1e+02 0.59461538 0.08083319
## 29  10.0 1e+02 0.59461538 0.08083319
## 30 100.0 1e+02 0.59461538 0.08083319
```

For polynomial, the best cost is 100, best degree is 2.


For radial, the best cost is 10, best gamma is 0.01.

(d)

Make some plots to back up your assertions in (b) and (c).


```
svm.linear = svm(mpglevel ~ ., data = Auto, kernel = "linear", cost = 1)
svm.poly = svm(mpglevel ~ ., data = Auto, kernel = "polynomial", cost = 100, degree = 2)
svm.radial = svm(mpglevel ~ ., data = Auto, kernel = "radial", cost = 10, gamma = 0.01)
plotsvm = function(fit) {
  for (name in names(Auto)[!(names(Auto) %in% c("mpg", "mpglevel", "name"))]) {
    plot(fit, Auto, as.formula(paste("mpg~", name, sep = "")))
  }
}
plotsvm(svm.linear)
```

SVM classification plot






Answer:




SVM classification plot



SVM classification plot

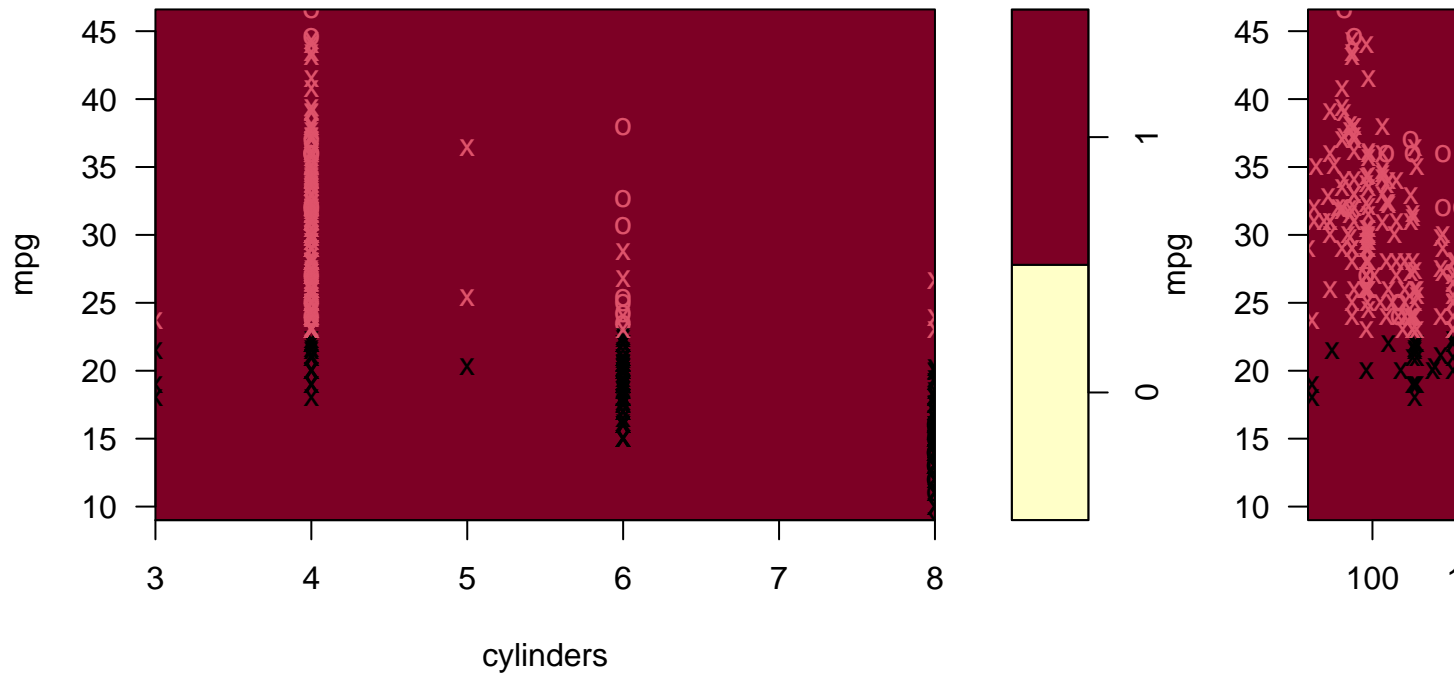


SVM classification plot




```
plotsvm(svm.poly)
```


SVM classification plot




SVM classification plot



SVM classification plot




SVM classification plot




```
plotsvm(svm.radial)
```


SVM classification plot




SVM classification plot



SVM classification plot



SVM classification plot



9.8

(a)

Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
set.seed(1)
train <- sample(dim(OJ)[1], 800)
OJ.train <- OJ[train, ]
OJ.test <- OJ[-train, ]
```

Answer:

(b)

Fit a support vector classifier to the training data using $\text{cost} = 0.01$, with Purchase as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics, and describe the results obtained.

```
svm.linear = svm(Purchase ~ ., kernel = "linear", data = OJ.train, cost = 0.01)
summary(svm.linear)
```

Answer:

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "linear", cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##         cost: 0.01
##
## Number of Support Vectors: 435
##
## ( 219 216 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

(c)

What are the training and test error rates?

```
train.pred = predict(svm.linear, OJ.train)
table(OJ.train$Purchase, train.pred)
```

Answer:

```
##      train.pred
##      CH  MM
## CH 420  65
## MM  75 240
```

```
(train.errorRate = (75+65)/(75+65+420+240))
```

```
## [1] 0.175
```

```
test.pred = predict(svm.linear, OJ.test)
table(OJ.test$Purchase, test.pred)
```

```
##      test.pred
##      CH  MM
## CH 153  15
## MM   33  69
```

```
(test.errorRate = (15+33)/(15+33+153+69))
```

```
## [1] 0.1777778
```

(d)

Use the tune() function to select an optimal cost. Consider values in the range 0.01 to 10.

```
set.seed(2)
```

```
tune.linear = tune(svm, Purchase ~ ., data = OJ.train, kernel = "linear", ranges = list(cost = seq(0.01, 10, length = 100)))
tune.linear$best.parameters$cost
```

Answer:

```
## [1] 4.01
```

(e)

Compute the training and test error rates using this new value for cost.

```
svm.linear2 = svm(Purchase ~ ., kernel = "linear", data = OJ.train, cost = tune.linear$best.parameters$cost)
```

```
train.pred2 = predict(svm.linear2, OJ.train)
table(OJ.train$Purchase, train.pred2)
```

Answer:

```
##      train.pred2
##      CH  MM
## CH 423  62
## MM   70 245
```

```
(train.errorRate2 = (61+70)/(61+70+424+245))
```

```
## [1] 0.16375
```

```
test.pred2 = predict(svm.linear2, OJ.test)
table(OJ.test$Purchase, test.pred2)
```

```
##      test.pred2
##      CH  MM
## CH 155  13
## MM   29  73
```

```
(test.errorRate2 = (13+29)/(13+29+155+73))
```

```
## [1] 0.1555556
```

(f)

Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for gamma.

```
set.seed(3)
tune.radial = tune(svm, Purchase ~ ., data = OJ.train, kernel = "radial", ranges = list(cost = seq(0.01, 10, length = 10)))
svm.radial = svm(Purchase ~ ., data = OJ.train, kernel = "radial", cost = tune.radial$best.parameters$cost)
train.pred.radial = predict(svm.radial, OJ.train)
table(OJ.train$Purchase, train.pred.radial)
```

Answer:

```
##      train.pred.radial
##      CH  MM
## CH 438  47
## MM  72 243
```

```
(47+72)/dim(OJ.train)[1]
```

```
## [1] 0.14875
```

```
test.pred.radial = predict(svm.radial, OJ.test)
table(OJ.test$Purchase, test.pred.radial)
```

```
##      test.pred.radial
##      CH  MM
## CH 150  18
## MM  30  72
```

```
(30+18)/dim(OJ.test)[1]
```

```
## [1] 0.1777778
```

(g)

Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set degree = 2.

```
set.seed(4)
tune.poly = tune(svm, Purchase ~ ., data = OJ.train, kernel = "poly", ranges = list(cost = seq(0.01, 10, length = 10)))
svm.poly = svm(Purchase ~ ., data = OJ.train, kernel = "poly", cost = tune.poly$best.parameters$cost)
train.pred.poly = predict(svm.poly, OJ.train)
table(OJ.train$Purchase, train.pred.poly)
```

Answer:

```
##      train.pred.poly
##      CH  MM
## CH 454  31
## MM  97 218
```

```
(31+97)/dim(OJ.train)[1]
```

```
## [1] 0.16
test.pred.poly = predict(svm.poly, OJ.test)
table(OJ.test$Purchase, test.pred.poly)
```

```
##      test.pred.poly
##           CH  MM
##    CH 154  14
##    MM  47  55
(47+14)/dim(OJ.test)[1]
```

```
## [1] 0.2259259
```

(h)

Overall, which approach seems to give the best results on this data?

Answer: Linear kernel has lowest test error rate.